

Dataset Description

We benchmarked on the **Bitcoin Alpha trust network**

- **Nodes:** 3,783
- **Edges:** 24,186 (directed, weighted, signed)
- **Edge weights:** -10 to +10 (trust score)
- **Source:** <https://snap.stanford.edu/data/soc-sign-bitcoin-alpha.html>

Member 1

1. Machine Specification

Component	Specification
Model	HP Pavilion x360
Name	
Chip	Intel core i5 10th Generation
Memory	8 GB
Storage	512 GB SSD
OS	Windows

2. Algorithms and Complexity Analysis

2.1 Dijkstra's Algorithm

Procedure Dijkstra(G, S)

BEGIN

Initialize distances: $\text{dist}[S] \leftarrow 0$, others $\leftarrow \infty$

Initialize priority queue PQ with (0, S)

WHILE PQ is NOT EMPTY **DO**

$(d, U) \leftarrow \text{EXTRACT-MIN}(\text{PQ})$

IF $d > \text{dist}[U]$ **THEN** CONTINUE

FOR EACH V in Neighbors of U **DO**

IF $\text{dist}[U] + \text{weight}(U, V) < \text{dist}[V]$ **THEN**

$\text{dist}[V] \leftarrow \text{dist}[U] + \text{weight}(U, V)$

$\text{parent}[V] \leftarrow U$

INSERT/UPDATE PQ with $(\text{dist}[V], V)$

END IF

END FOR

END WHILE

END

Best Case:

Time: $O((V + E) \log V)$

- Occurs when the graph is sparse ($E \approx V$) and the priority queue operations are efficient. Each vertex is extracted once ($O(V \log V)$), and each edge is relaxed once ($O(E \log V)$).
- In our dataset ($\sim 3,783$ nodes, $\sim 15,000$ edges), this is typical since E is relatively small compared to V^2 .

Average Case:

Time: $O((V + E) \log V)$

- On average, Dijkstra's explores most vertices and edges, with priority queue operations dominating the runtime. The binary heap implementation ensures $O(\log V)$ per operation (extract-min and decrease-key).

Worst Case:

Time: $O((V + E) \log V)$

- Happens when the graph is dense ($E \approx V^2$). Each vertex is processed once, and each edge relaxation involves a priority queue update. For our dataset, this isn't the case, but the complexity remains the same due to the priority queue.

Space Complexity:

- $O(V)$ for the distance array, parent array, and priority queue (stores up to V vertices).
- Total space: $O(V)$.

2.2 Bellman-Ford Algorithm

Procedure BellmanFord(G, S)

BEGIN

Initialize distances: $\text{dist}[S] \leftarrow 0$, others $\leftarrow \infty$

FOR $I = 1$ TO $|V| - 1$ DO

 FOR EACH edge (U, V, w) in $G.E$ DO

 IF $\text{dist}[U] + w < \text{dist}[V]$ THEN

$\text{dist}[V] \leftarrow \text{dist}[U] + w$

$\text{parent}[V] \leftarrow U$

 END IF

 END FOR

 IF no updates in this iteration THEN BREAK

END FOR

END

Best Case:

Time: $O(V + E)$

- Occurs when early termination kicks in after the first iteration (e.g., a graph where the source is isolated or distances stabilize quickly). Only one pass over all edges is needed.

Average Case:

Time: $O(VE)$

- Typically, Bellman-Ford performs $V-1$ iterations over all edges to ensure shortest paths are found. For our dataset ($\sim 3,783$ nodes, $\sim 15,000$ edges), this is significant but mitigated by early termination if updates cease.

Worst Case:

Time: $O(VE)$

- Happens when the graph requires the full $V-1$ iterations to compute shortest paths (e.g., a path graph where distances update in each iteration). For our dataset, this means $\sim 3,783 \times 15,000$ operations, which is computationally expensive.

Space Complexity:

- $O(V)$ for the distance and parent arrays.
- $O(E)$ for storing the edge list (though this is part of the input).
- Total space: $O(V)$, excluding the edge list.

2.3 Diameter (Using All-Pairs Dijkstra's)

Procedure ComputeDiameter(G)

BEGIN

Initialize $\text{max_dist} \leftarrow 0$

FOR EACH U in $G.V$ DO

 CALL Dijkstra(G, U) to get distances $\text{dist}[]$

 FOR EACH V in $G.V$ DO

```

    IF dist[V]  $\neq \infty$  AND dist[V] > max_dist THEN

        max_dist  $\leftarrow$  dist[V]

        Update longest path

    END IF

END FOR

END FOR

RETURN max_dist

END

```

Best Case:

Time: $O(V(V + E) \log V)$

- Occurs when the graph is sparse ($E \approx V$), and Dijkstra's is efficient. Each Dijkstra's run takes $O((V + E) \log V)$, and we run it V times. For our dataset ($\sim 3,783$ nodes, $\sim 15,000$ edges), this is the typical case.

Average Case:

Time: $O(V(V + E) \log V)$

- On average, each Dijkstra's call explores most vertices and edges, multiplied by V calls. The priority queue operations dominate, as in Dijkstra's.

Worst Case:

Time: $O(V(V + E) \log V)$


- Happens when the graph is dense ($E \approx V^2$), making each Dijkstra's run $O((V + V^2) \log V) = O(V^2 \log V)$. Total runtime becomes $O(V^3 \log V)$. Our dataset is sparse, so we avoid this worst case.

Space Complexity:


- $O(V)$ for the distance and parent arrays in each Dijkstra's call.
- $O(V)$ for the priority queue per call.
- Total space: $O(V)$, as arrays are reused across calls.

3. Algorithm Graph Plots


3.1 Dijkstra Performance:



3.2 Bellmanford Performance:



3.3 Diameter Performance:



4. Implementation Details

4.1 Data Structures

Adjacency:

- **Dijkstra's, Bellman-Ford, and Diameter:** Adjacency list implemented as a vector<vector<Edge>> for Dijkstra's and Diameter, and a vector<Edge> for Bellman-Ford. Each Edge struct stores (to, weight) for Dijkstra's and Diameter, and (from, to, weight) for Bellman-Ford, reflecting the directed graph from the dataset (source, target, weight, timestamp).
- **BFS (Not Implemented):** Would use a queue<int> for node indices, with $O(1)$ amortized enqueue/dequeue.
- **DFS & Cycle Detection (Not Implemented):** Would use recursion with a call stack, supplemented by bool Visited[] and bool InStack[] arrays for cycle detection.

Impact on Complexity:

- Core operations (vector access, queue operations, edge traversal) are $O(1)$ per element.
- For Dijkstra's, priority queue (implemented via priority_queue) adds $O(\log V)$ per insert/update due to heap operations.
- Overall algorithmic complexity ($O((V + E) \log V)$ for Dijkstra's, $O(VE)$ for Bellman-Ford, $O(V(V + E) \log V)$ for Diameter) is governed by graph size and algorithm logic, not data structure overhead beyond the priority queue in Dijkstra's.

4.2 Choice of Stack vs. Queue

- **Dijkstra's:**
 - Uses a priority queue (priority_queue<pair<long long, int>, vector<pair<long long, int>>, greater<>>) instead of a standard queue or stack. This choice ensures the minimum distance node is always processed next, critical for greedy shortest-path computation. A queue (FIFO) or stack (LIFO) would not maintain the order needed for optimal path selection.

- **Bellman-Ford:**

- Does not use a queue or stack explicitly; relies on iterative edge relaxation over all edges. The absence of a dynamic data structure simplifies implementation but requires $V-1$ passes, leveraging the algorithm's inherent structure rather than stack/queue ordering.

- **Diameter:**

- Internally uses Dijkstra's with a priority queue for each source node to compute all-pairs shortest paths. The priority queue choice mirrors Dijkstra's, ensuring efficient extraction of minimum distance nodes. A queue or stack would not suit the all-pairs approach, as it requires repeated shortest-path computations from each vertex.

Member 2

1. Machine Specification

Component	Specification
Chip	Intel core i5 6th Generation
Memory	8 GB
Storage	512 GB SSD
OS	Windows

2. Algorithms and Complexity Analysis

2.1 Prim's Algorithm (MST)

Procedure PrimMst.run(Graph, Start Node)

BEGIN

 Initialize an empty set visited

 Initialize a priority queue pq (min-heap)

 Mark start as visited

 Push all edges from start to pq

 WHILE pq is NOT EMPTY DO

 Pop the minimum weight edge (u, v)

 IF v is NOT visited THEN

 Add (u, v) to MST

 Mark v as visited

 Push all edges from v to pq

 END IF

 END WHILE

END

Best Case:

- **Time: $O(E \log V)$**
- When graph is sparse and priority queue operations dominate. Each edge is added to the queue at the most once.

Average Case:

- **Time: $O(E \log V)$**
- On average performance remains $O(E \log V)$ as heap operations determine performance in typical cases.

Worst Case:

- **Time: $O(E \log V)$**
- All edges are pushed into the priority queue, and we perform heap operations for each, resulting in $O(E \log V)$.

Space Complexity:

- $O(V)$ for the visited array and $O(E)$ for the priority queue.

2.2 Kruskal Algorithm (MST)

Procedure KruskalMST.run(Graph)

BEGIN

 Initialize disjoint sets using Union-Find

 Extract all edges and sort by weight

 FOR each edge (u, v) in ascending order of weight DO

 IF u and v are in different components THEN

 Add edge to MST

 Union their components

 END IF

 END FOR

END

Time Complexity:

- **Best, Average, and Worst Case: $O(E \log E)$**
 - Best Case: Sorting edges dominates the runtime; few union operations are needed.
 - On average, sorting still dominates, and Union-Find operations are near constant with path compression.
 - Worst: All edges need sorting and checked for cycles.

Space Complexity:

- $O(V)$ for Union-Find parent and rank maps
- $O(E)$ for storing edges
- Total: $O(V + E)$

2.3 Average Degree Calculation

Time Complexity:

- **Best, Average, and Worst Case: $O(E)$**
 - Best: File read is linear in the number of edges.
 - Average: Each edge is processed once with constant-time map and set operations.
 - Worst: Each edge is processed once with constant-time map and set operations.

Space Complexity:


- $O(V)$ for degree map and nodes set
- So, the total space complexity = $O(V)$

Average degree txt file


```
degree > ≡ average_degree.txt
1  Average Degree: 12.7867
2  Total Nodes: 3783
3  Total Edges: 24186
4  
```

3. Algorithm Graph Plots


3.1 Prims Algorithm



3.2 Kruskal Algorithm



3.3 Average Degree



4. Implementation Details

4.1 Data Structures

- **Graph:** Stored as an adjacency list using `unordered_map<int, vector<pair<int, double>>>`
- **Prim's:** Uses `set<int>` for visited nodes and a **min-heap** (priority queue) to select the smallest edge.
- **Kruskal's:** Uses a **sorted edge list** and a **Union-Find** structure (`unordered_map` for parent and rank).
- **Average Degree:** Uses `unordered_map<int, int>` for node degrees and `set<int>` for unique nodes.
- All operations like insertion, access, and updates are $O(1)$ or $O(\log E)$, so they don't impact overall complexity.

4.2 Choice of Stack vs. Queue

- Prim's: Uses a min-heap (not a stack/queue) to ensure greedy selection.
- Kruskal's: Relies on Union-Find, no queue or stack involved.
- Average Degree: Processes data linearly—no need for stack/queue.
- The selected structures ensure efficiency; using incorrect ones (e.g., a regular queue in Prim's) would break correctness or increase time complexity.

4.3 Execution Time

- **Prim Algo : 113300 ms**
- **Kruskal Algo: 104974 ms**
- **Average Degree: 206130 ms**

Member 3

1. Machine Specification

Component	Specification
Model Name	MacBook Pro (16-inch, Mac16,8)
Chip	Apple M4 Pro (8 performance + 4 efficiency cores)
Memory	24 GB LPDDR5
Storage	500 GB SSD
OS	macOS 15.4.1

2. Algorithms and Complexity Analysis

2.1 Breadth-First Search (BFS)

Procedure BFS(G, S)

 BEGIN

 FOR I = 0 TO Size of G.V - 1 STEP 1 DO

 Visited[I] \leftarrow FALSE

 END FOR

 Create an empty queue Q

 Visited[S] \leftarrow TRUE

 ENQUEUE(Q, S)

 WHILE Q is NOT EMPTY DO

 U \leftarrow DEQUEUE(Q)

 FOR EACH V in Neighbors of U DO

 IF Visited[V] = FALSE THEN

 Visited[V] \leftarrow TRUE

 ENQUEUE(Q, V)

 END IF

 END FOR

 END WHILE

 END

Best Case:

- **Time: $O(V+E)$**
- This is $O(V+E)$ because BFS must check every vertex and edge at most once.

Average Case:

- **Time: $O(V+E)$**
- On average, it explores all vertices and their edges in the connected component of the starting node.

Worst Case:

- **Time: $O(V+E)$**
- Happens when the entire graph is connected, so all vertices and all edges are explored.

Space Complexity:

- $O(V)$ for the visited array and the queue in the worst case (when all vertices are enqueued).

2.2 Depth-First Search (DFS)

Procedure DFS(G, U)

BEGIN

Visited[U] \leftarrow TRUE

FOR EACH V in Neighbors of U DO

IF Visited[V] = FALSE THEN

DFS(G, V)

END IF

END FOR

END

Time Complexity:

- **Best, Average, and Worst Case: $O(V+E)$**

- Every vertex is visited once $\rightarrow O(V)$
- Every edge is explored once $\rightarrow O(E)$
- So, total = $O(V+E)$

Space Complexity:

- **Visited array:** $O(V)$ to track visited nodes

2.3 Cycle Detection in Directed Graphs

Procedure DetectCycle(G)

BEGIN

```
    FOR EACH V in G.V DO
        Visited[V] ← FALSE
        InStack[V] ← FALSE
    END FOR
```

```
    FOR EACH V in G.V DO
        IF Visited[V] = FALSE THEN
            IF CALL DFS_Cycle(V) THEN
                RETURN TRUE
            END IF
        END IF
    END FOR
```

```
    RETURN FALSE
```

END

Procedure DFS_Cycle(U)

BEGIN

```
    Visited[U] ← TRUE
    InStack[U] ← TRUE
```

```
    FOR EACH V in Neighbors of U DO
        IF InStack[V] = TRUE THEN
            RETURN TRUE // Back edge ⇒ cycle
        ELSE IF Visited[V] = FALSE AND CALL DFS_Cycle(V) THEN
            RETURN TRUE
        END IF
    END FOR
```

```
InStack[U] ← FALSE  
RETURN FALSE
```

END

Time Complexity:

- **Best, Average, and Worst Case: $O(V+E)$**
 - Every vertex is visited once, and every edge is explored once during the DFS traversal. The cycle check involves checking neighbors but still follows the same traversal pattern.

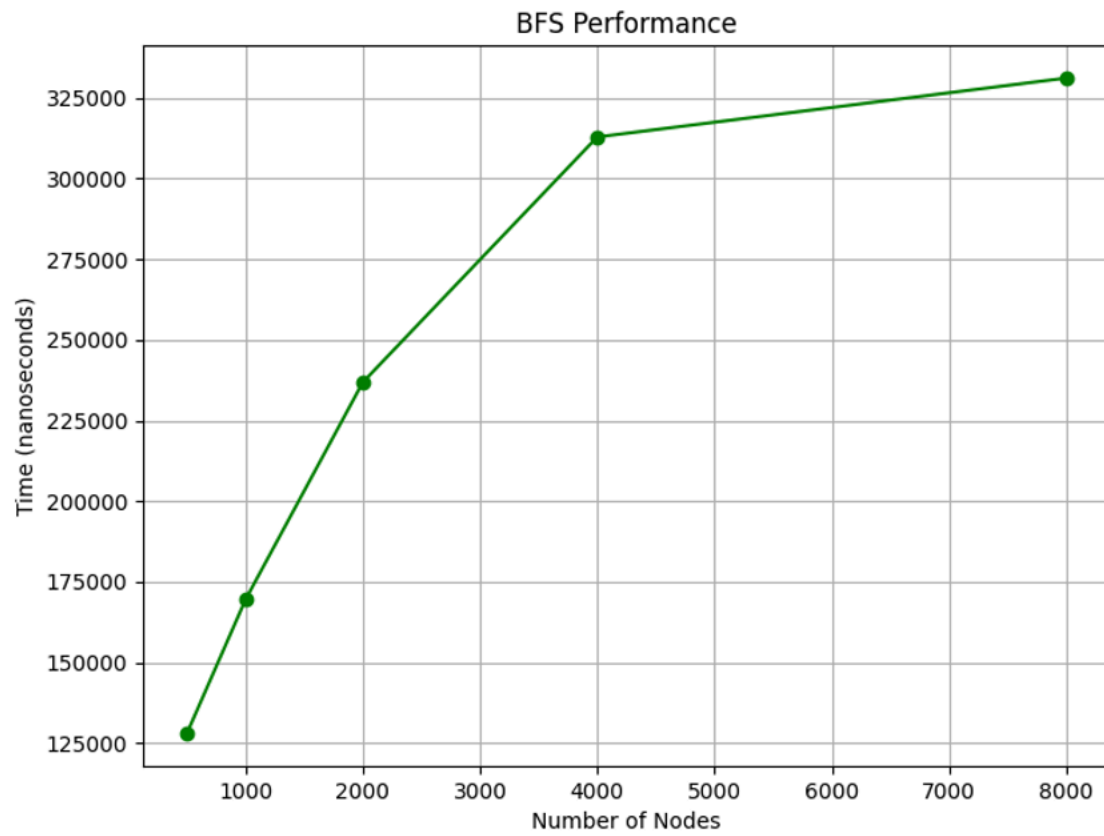
Space Complexity:

- Visited array: $O(V)$
- InStack array: $O(V)$
- Call stack (in DFS): $O(V)$
- So, the total space complexity = $O(V)$

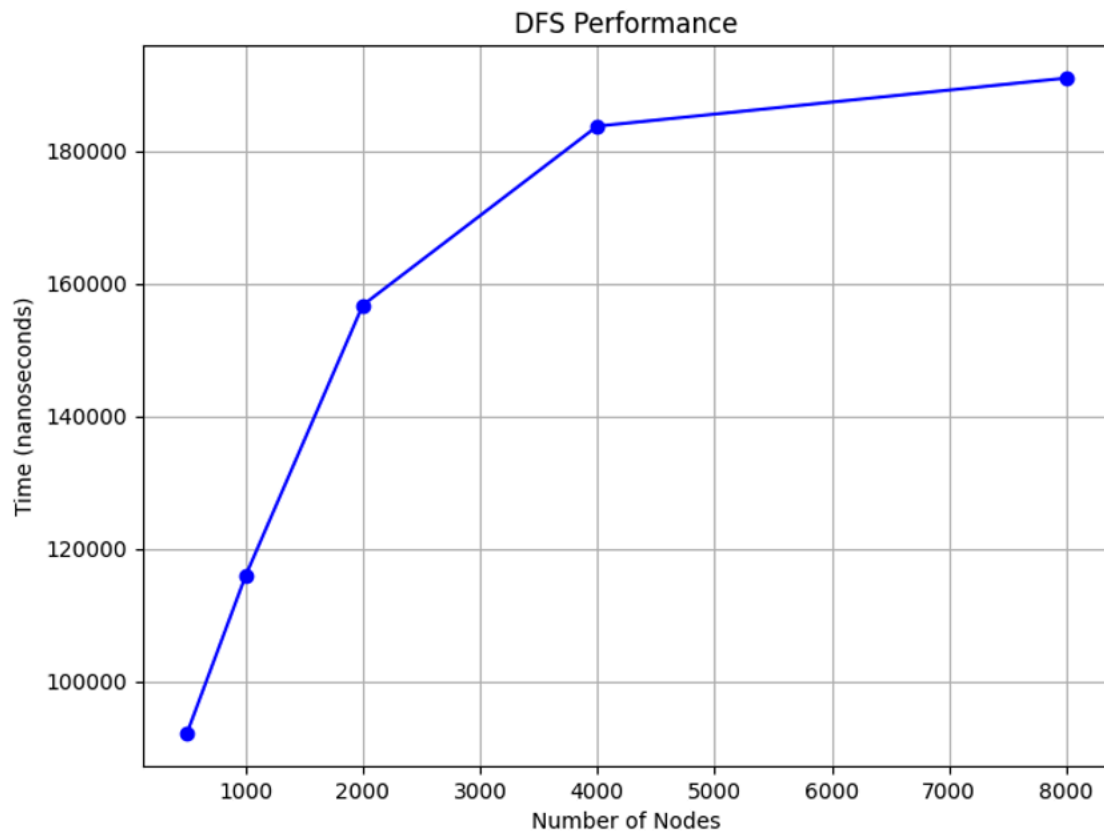
3. Algorithm Graph Plots

3.1 BFS Performance

Time grows roughly linearly with $|V| + |E|$.



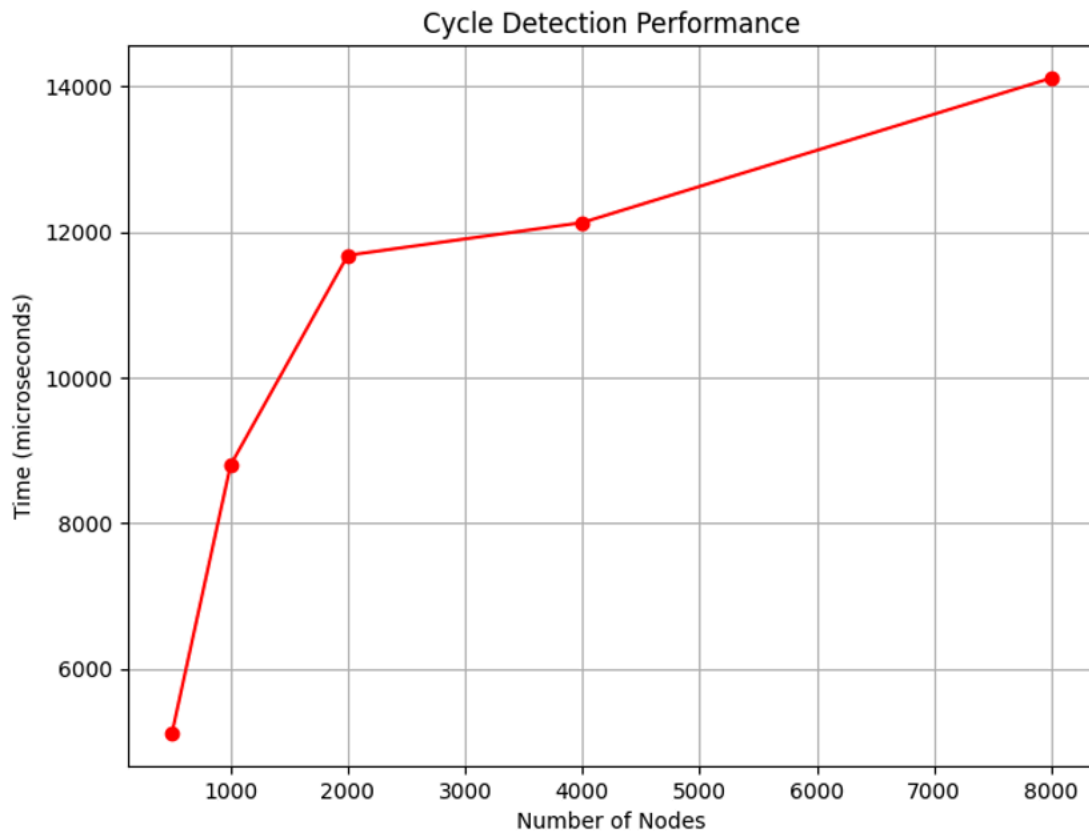
3.2 DFS Performance



Similar linear trend; slightly faster than BFS

3.3 Cycle Detection Performance

Measured in microseconds, same in linear $O(V+E)$.



4. Implementation Details

4.1 Data Structures

Adjacency:

- LinkedList per vertex for neighbors; each node stores (target, rating, timestamp).

BFS:

- Queue built from the same Node type.
- Enqueue/dequeue in $O(1)$ amortized.

DFS & Cycle Detection:

- Recursive calls manage their own call-stack.

- Additionally for cycle detection, we maintain a separate boolean `inRecursionStack[]`.

Impact on Complexity:

- All core operations (enqueue, dequeue, push, pop, neighbor traversal) remain $O(1)$.
- Thus overall algorithmic complexity is governed purely by the graph size, not by data-structure overhead.

4.2 Choice of Stack vs. Queue

· DFS via Recursion (Call Stack)

I deliberately used recursion instead of a Stack class. With recursion, when you visit a node and push all its neighbors onto the call stack, the *first* neighbor pushed is automatically the *first* one you explore, exactly matching DFS order. If I use a stack, I must reverse the neighbor list each time to achieve the same behavior.

· BFS via Queue

For breadth-first search, a FIFO Queue is the natural and only data structure that can be used.

5. Graph Visualization

